Final Team Project

by Tommy Barron, April Chia, and Taylor Kirk

Data Importing and Pre-processing

Import data set and describe characteristics

- Dataset 3
- Command: df <- read_csv("online_shoppers_intention.csv")

dim(df) str(df)

```
Γ17 12330
spc_tbl_[12,330 \times 18] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
$ Administrative
                         : num [1:12330] 0 0 0 0 0 0 0 1 0 0 ...
 $ Administrative_Duration: num [1:12330] 0 0 0 0 0 0 0 0 0 0 ...
$ Informational
                          : num [1:12330] 0 0 0 0 0 0 0 0 0 0 ...
 $ Informational_Duration : num [1:12330] 0 0 0 0 0 0 0 0 0 0 ...
$ ProductRelated
                          : num [1:12330] 1 2 1 2 10 19 1 0 2 3 ...
 $ ProductRelated_Duration: num [1:12330] 0 64 0 2.67 627.5 ...
 $ BounceRates
                          : num [1:12330] 0.2 0 0.2 0.05 0.02 ...
 $ ExitRates
                         : num [1:12330] 0.2 0.1 0.2 0.14 0.05 ...
 $ PageValues
                         : num [1:12330] 0 0 0 0 0 0 0 0 0 0 ...
 $ SpecialDay
                          : num [1:12330] 0 0 0 0 0 0 0.4 0 0.8 0.4 ...
 $ Month
                          : chr [1:12330] "Feb" "Feb" "Feb" "Feb" ...
 $ OperatingSystems
                          : num [1:12330] 1 2 4 3 3 2 2 1 2 2 ...
                          : num [1:12330] 1 2 1 2 3 2 4 2 2 4 ...
 $ Browser
 $ Region
                          : num [1:12330] 1 1 9 2 1 1 3 1 2 1 ...
 $ TrafficType
                          : num [1:12330] 1 2 3 4 4 3 3 5 3 2 ...
$ VisitorType
                          : chr [1:12330] "Returning_Visitor" "Returning_Visitor"
"Returning_Visitor" "Returning_Visitor" ...
 $ Weekend
                          : logi [1:12330] FALSE FALSE FALSE FALSE TRUE FALSE ...
                          : logi [1:12330] FALSE FALSE FALSE FALSE FALSE ...
 $ Revenue
```

Duplicates

Command: df_dup <- df[duplicated(df),]
 dim(df_dup)

```
[1] 121 18
```

- Command: df <- df[!duplicated(df),]
- **Purpose:** To remove duplicate rows

Finding missing values

Command: na_counts <- sapply(df, function(x) sum(is.na(x))) na_counts

Informational_Duration	Informational	Administrative_Duration	Administrative
0	128	0	0
ExitRates	BounceRates	ProductRelated_Duration	${\tt ProductRelated}$
0	0	0	0
OperatingSystems	Month	SpecialDay	PageValues
123	0	0	135
VisitorType	TrafficType	Region	Browser
0	0	0	0
		Revenue	Weekend
		0	0

Outliers

- Viewed boxplots
- **Result**: left outliers in the data set

Handling missing values (Informational)

```
# Informational column without the 0's
info_nzero <- subset.data.frame(df[df$Informational > 0, ])
# Probability table for unique values
probability <- table(info_nzero$Informational, useNA = "no")</pre>
prob_table <- prop.table(probability)</pre>
# Function to randomly impute one of those values based on their proportion
random_impute <- function(values, prob_table, size) {</pre>
  sample(values, size, replace = T, prob = prob_table)
# Separating the NA Information rows from the ones where InformationDuration is greater than 0, and the ones
where it's less than or equal to 0
na_indices <- which(is.na(df$Informational))</pre>
na_indices_great_zero <- na_indices[df[na_indices, "Informational_Duration"] > 0]
na_indices_zero_or_less <- na_indices[df[na_indices, "Informational_Duration"] <= 0]</pre>
# Converting the probability table of unique values to numeric
unique_values <- as.numeric(names(prob_table))</pre>
probs <- as.numeric(prob_table)</pre>
```

Command: df\$Informational[na_indices_great_zero] <- random_impute(unique_values, probs, length(na_indices_great_zero))

df\$Informational[is.na(df\$Informational)] <- 0

Handling missing values (PageValues)

```
# Subsetting the PageValues NA values
pv_na <- df[!is.na(df$PageValues), ]

# Separate PageValues if Administrative, Information, or ProductRelated columns are greater than 0

pv_na <- pv_na[pv_na$Administrative > 0 | pv_na$Informational > 0 | pv_na$ProductRelated > 0, ]
summary(pv_clean$PageValues)
```

```
Min. 1st Qu. Median Mean 3rd Qu. Max.
0.000 0.000 0.000 5.973 0.000 361.764
```

Command: df\$PageValues[is.na(df\$PageValues)] <- 0

Handling missing values (OperatingSystems)

123 missing values

Command: df <- df[!is.na(df\$OperatingSystems),]

Purpose: To remove rows with missing values

Transformation of the data

- Recoding the Special Day column to a categorical
 - Command:


- Converting Month column to factor
 - Command: df\$Month_factor <- as.factor(df\$Month)
- Converting BounceRates and ExitRates into percentage values
 - Command:


```
df_percent <- df
df_percent$BounceRates <- df_percent$BounceRates * 100
df_percent$ExitRates <- df_percent$ExitRates * 100</pre>
```

Cleaned dataframe


Command: data <- df_percent


Data Analysis and Visualization








Correlation Plot
















Proportion Table for Month Category


	880	
	False	True
Aug	0.82201405	0.17798595
Dec	0.86797066	0.13202934
Feb	0.98136646	0.01863354
Jul	0.84210526	0.15789474
June	0.89259259	0.10740741
Mar	0.89164786	0.10835214
May	0.88543628	0.11456372
Nov	0.73500697	0.26499303
Oct	0.79014599	0.20985401
Sep	0.80630631	0.19369369






Proportion Table for Special Day Category

False	True
0.82603325	0.17396675
0.91764706	0.08235294
0.94420601	0.05579399
0.91369048	0.08630952
0.96507937	0.03492063
0.93377483	0.06622517
	0.82603325 0.91764706 0.94420601 0.91369048 0.96507937





Data Analytics

Preliminary Logistic Regression w/ all Variables

```
# Subset and clean up data
reg.data <- subset(df, select=c(1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18))
reg.data$Revenue <- as.logical(reg.data$Revenue)
reg.dataSRevenue <- as.numeric(reg.dataSRevenue)
# Create Train and Test data sets
trainIndex <- createDataPartition(reg.data$Revenue,
                               p = .8,
                               list = FALSE,
                               times = 1)
train <- reg.data[trainIndex,]
test <- req.data[-trainIndex,]
# Logistic Regression Model
model <- qlm(Revenue~., family = binomial(link="logit"), data=train)
summary(model)
Administrative
                              1.123e-03 1.262e-02
                                                     0.089 0.92912
Administrative Duration
                             -1.317e-04 2.267e-04
                                                    -0.581 0.56139
Informational
                              4.954e-02 2.992e-02
                                                    1.656 0.09781 .
Informational_Duration
                             -2.706e-06 2.444e-04 -0.011 0.99117
ProductRelated
                              5.203e-04 1.401e-03
                                                    0.371 0.71040
ProductRelated_Duration
                              9.653e-05 3.375e-05
                                                    2.860 0.00424 **
BounceRates
                             -4.901e-02 3.795e-02
                                                    -1.291
ExitRates
                             -1.557e-01 2.708e-02 -5.750 8.92e-09 ***
PageValues
                              8.178e-02 2.747e-03 29.768
SpecialDay
                             -1.136e-01 2.733e-01
                                                    -0.416
MonthDec
                             -5.457e-01 2.108e-01
                                                    -2.589 0.00963 **
MonthFeb
                             -1.555e+00 6.531e-01
MonthJul
                              1.444e-01 2.499e-01
MonthJune
                             -2.972e-01 3.091e-01
                                                    -0.961 0.33635
MonthMar
                             -4.807e-01 2.082e-01 -2.309 0.02094 *
MonthMay
                             -5.839e-01 2.024e-01
                                                    -2.885 0.00392 **
MonthNov
                              6.115e-01 1.891e-01
                                                            0.00122 **
Monthoct
                             -1.491e-01 2.377e-01 -0.627
                                                            0.53060
                              1.062e-01 2.400e-01
                                                     0.443 0.65810
MonthSep
OperatingSystems
                             -1.117e-01 4.416e-02
                                                    -2.529
                                                            0.01144
Browser
                              4.260e-02 2.106e-02
                                                     2.023
                                                            0.04307
Region
                             -1.547e-02 1.487e-02
                                                    -1.040 0.29821
TrafficType
                              1.010e-02 9.216e-03
                                                    1.096
                                                            0.27302
VisitorTypeOther
                             -3.815e-01 5.985e-01
                                                    -0.637 0.52385
VisitorTypeReturning_Visitor -3.083e-01 9.728e-02
                                                    -3.170
                                                            0.00153 **
WeekendTRUE
                              8.571e-02 8.040e-02
                                                    1.066 0.28642
```

- Most significant predictor variables:
 - ProductRelated_Duration
 - ExitRates
 - PageValues
 - Month
 - VisitorType

Final Logistic Regression Model

```
# Set Seed
set.seed(123)
# Create Train and Test data sets
sample <- sample(c(TRUE, FALSE), nrow(df), replace=TRUE, prob=c(0.8,0.2))</pre>
train <- df[sample,
test <- df[!sample, ]
# Logistic Regression Model
model <- olm(Revenue~ProductRelated Duration+ExitRates+PageValues+MonthFeb+MonthMar+Month
May+MonthNov+MonthDec+ReturningVisitor, family="binomial", data=train)
call:
glm(formula = Revenue ~ ProductRelated_Duration + ExitRates +
   PageValues + MonthFeb + MonthMar + MonthMay + MonthNov +
   MonthDec + Returning Visitor, family = "binomial", data = train)
coefficients:
                       Estimate Std. Error z value
                                                              Pr (> |z|)
(Intercept)
                                  ProductRelated Duration 0.20939
                                  0.02988 7.009
                                                       0.0000000000024 ***
                                 0.08812 -10.026 < 0.00000000000000000 ***
ExitRates
                       -0.88345
Pagevalues
                       1.51171
                                  MonthFeb
                       -1.63946
                                 0.63078 -2.599
                                                              0.009347 **
                       -0.58883
                                  0.13123 -4.487
                                                       0.0000072208784 ***
MonthMar
                       -0.52410
                                  0.10907 -4.805
                                                       0.0000015459376 ***
MonthMay
MonthNov
                       0.53690
                                  0.09624 5.579
                                                       0.0000000242312 ***
Month Dec
                       -0.47760
                                  0.12864 -3.713
                                                             0.000205 ***
ReturningVisitor
                       -0.30872
                                  0.09330 -3.309
                                                             0.000937 ***
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
(Dispersion parameter for binomial family taken to be 1)
   Null deviance: 8496.8 on 9808 degrees of freedom
Residual deviance: 5776.9 on 9799 degrees of freedom
AIC: 5796.9
Number of Fisher Scoring iterations: 7
```

- Response Variable:
 - Revenue
- Predictor Variables:
 - Product Related Duration
 - Exit Rates
 - Page Values
 - February, March, May, November,
 December
 - Returning Visitor
- Create a Train and Test using sample()
 - Train = 80%
 - o Test = 20%

Evaluation of Model Part 1/2

```
# Predicted results and Accuracy
predicted <- predict(model, test, type="response")
predicted_results <- ifelse(predicted > 0.28, 1, 0)

misclasificError <- mean(predicted_results != test$Revenue)
print(paste('Accuracy', 1-misclasificError))

# Calculation of pseudo R-Squared Value
# pscl comes from library(pscl)
pscl::pR2(model)["McFadden"]

[1] "Accuracy 0.894166666666667"
fitting null model for pseudo-r2
    McFadden
0.3201154</pre>
```

- Test the model using the Test data
- R² values apply differently to logistic regression
- McFadden Pseudo r-squared value for Logistic Regression
 - "DL McFadden stated that a pseudo-R² higher than 0.2 represents an excellent fit"
 - One-Off Coder (2024)

Evaluation of Model Part 2/2

```
# Check the predictor variables for multicollinearity
car::vif(model)
# Confusion Matrix to check false positive and false negatives
conf_matrix <- table(Predicted = predicted_results, Actual = test$Revenue)</pre>
cat(sprintf("Confusion Matrix:\n"))
conf matrix
ProductRelated Duration
                                     ExitRates
                                                           PageValues
              1.119017
                                      1.095266
                                                             1.048322
              MonthFeb
                                      MonthMar
                                                             MonthMay
              1.013745
                                      1.373842
                                                             1.622092
              MonthNov
                                      MonthDec
                                                     ReturningVisitor
              1.836362
                                      1.390703
                                                             1.112886
Confusion Matrix:
              Actual
Predicted FALSE TRUE
                 1930
                           161
                           216
```

- Check for Multicollinearity
- Check confusion matrix for false positives and false negatives
 - False Positives: 93
 - False Negatives: 161

ROC Curve

```
# Receiver Operating Characteristic Curve
# roc() comes from library(pROC)
roc_curve <- roc(test$Revenue, predicted)
plot(roc_curve)
auc_value <- auc(roc_curve)</pre>
print(auc_value)
Sensitivity
0.4 0.6
  0.0
    15
              10
                         0.5
                                    0.0
                                              -0.5
                       Specificity
Setting levels: control = FALSE, case = TRUE
Setting direction: controls < cases
Area under the curve: 0.8945
```

- AUC score interpretation:
 - = 0.5 is akin to random guessing
 - o 0.5-0.7 is slight strength
 - 0.7-0.8 is moderate strength
 - o 0.8-0.9 is good strength
 - o > 0.9 is great strength
- Evidently AI team (n.d.)

PR Curve

```
# Precision-Recall Curve
# pr.curve() comes from library(PRROC)
pr_curve <- pr.curve(scores.class0 = predicted[test$Revenue == 1],</pre>
                      scores.class1 = predicted[test$Revenue == 0],
                      curve = T)
plot(pr_curve)
precision <- conf_matrix[2, 2] / sum(conf_matrix[2, ])</pre>
recall <- conf_matrix[2, 2] / sum(conf_matrix[, 2])</pre>
print(paste("Precision Score:", precision))
print(paste("Recall Score:", recall))
f1_score <- 2 * ((precision * recall) / (precision + recall))
print(paste("F1 Score:", f1_score))
                            PR curve
                        AUC = 0.6383396
                                                                    0.8
    00
    0
                                                                    9.0
    9.0
    0.4
                                                                    0.4
    N
                                                                    0.2
    0
                 0.2
                                   0.6
                                            0.8
                                                      10
        0.0
                              Recall
```

- F1 Score Interpretation:
 - o < 0.5 is Not Good
 - o 0.5 0.8 is Ok
 - o 0.8 0.9 is Good
 - > 0.9 is Very Good
- Encord Blog (2023)

- .] "Precision Score: 0.699029126213592"
- [1] "Recall Score: 0.572944297082228"
- [1] "F1 Score: 0.629737609329446"

References

Encord Blog. (2023, July 18) F1 Score in Machine Learning. Encord.

https://encord.com/blog/f1-score-in-machine-learning/#:~:text=The%20F1%20score%2 Oranges%20between%200%20and%201%2C,can%20concurrently%20attain%20high% 20precision%20and%20high%20recall.

- Evidently AI. (n.d.) How to explain the ROC curve and ROC AUC score. Evidently AI. https://www.evidentlyai.com/classification-metrics/explain-roc-curve
- One-Off Coder. (2024, Apr. 03). *Pseudo r-squared for logistic regression*. Data Science Topics. https://datascience.oneoffcoder.com/psuedo-r-squared-logistic-regression.html
- Shah, C. (2020). Hands-on Introduction to Data Science. Cambridge University Press. https://doi.org/10.1017/9781108560412